

Machine Learning Framework for Air Quality Prediction

Annapureddy Siva Meenakshi
Department of Electrical engineering
IIT Hyderabad
ee23btech11008

Abstract—Air quality prediction is crucial for public health and environmental monitoring. This paper presents a comprehensive machine learning framework for predicting Air Quality Index (AQI) and individual pollutant concentrations across multiple cities. I employ Multi-Layer Perceptron (MLP), Recurrent Neural Networks (RNN), Long Short-Term Memory (LSTM) networks, and XGBoost algorithms, enhanced with hyperparameter tuning and multi-city comparative analysis. My methodology incorporates temporal features, lag variables, and rolling statistics to capture time-series patterns. Experimental results demonstrate that XGBoost with hyperparameter tuning achieves superior performance with an MAE of 0.34 for pollutant prediction and 0.85 for AQI calculation. The study extends beyond single-city analysis to compare air quality patterns across six major global cities, providing insights into urban air pollution dynamics and model generalization capabilities. The proposed framework demonstrates robust predictive performance across diverse geographical locations, making it suitable for operational air quality forecasting systems.

Index Terms—Air Quality Prediction, Machine Learning, XGBoost, LSTM, Neural Networks, Environmental Monitoring, Time Series Analysis

I. INTRODUCTION

Air pollution poses significant threats to public health and environmental sustainability, with the World Health Organization estimating millions of premature deaths annually attributable to poor air quality [7]. The increasing urbanization and industrial activities have exacerbated air pollution problems, making accurate air quality prediction an essential tool for environmental management and public health protection. Accurate prediction of air quality indices enables proactive measures and informed decision-making for urban planning, industrial regulation, and public health interventions.

Traditional methods often struggle to capture the complex, non-linear relationships in environmental data, making advanced machine learning approaches particularly suitable for this domain. The temporal nature of air quality data, influenced by meteorological conditions, seasonal patterns, and human activities, requires sophisticated modeling techniques that can capture both short-term and long-term dependencies.

This research presents a comprehensive machine learning framework for multi-pollutant air quality prediction through a comparative study of neural networks and ensemble methods. My key contributions include: (1) development of an ensemble feature engineering approach incorporating temporal patterns and historical dependencies, (2) rigorous evaluation of mul-

multiple algorithms including Multi-Layer Perceptron, Recurrent Neural Networks, Long Short-Term Memory networks, and XGBoost, (3) implementation of hyperparameter optimization for enhanced model performance, and (4) extensive multi-city analysis to validate model generalization across diverse urban environments.

II. RELATED WORK

Machine learning approaches have shown significant promise in environmental monitoring and air quality prediction. Previous studies have explored various techniques for forecasting air pollution levels, with comparative analyses demonstrating the effectiveness of different algorithms for specific pollutants and geographical contexts.

Recent research has focused on comparing machine learning methods for predicting specific pollutants. Liu et al. [1] conducted a comprehensive comparison of machine learning methods for ground-level ozone prediction in Beijing, highlighting the varying performance of different algorithms across seasonal variations. Their work emphasized the importance of algorithm selection based on specific pollutant characteristics and environmental conditions.

Deep learning approaches, particularly LSTM networks, have demonstrated strong capabilities in capturing temporal dependencies in air quality data. Naresh and Indira [5] developed a multivariate LSTM model for air pollution prediction, showing improved performance through the integration of multiple environmental variables. Similarly, Blanco et al. [3] proposed a machine learning approach leveraging satellite observations and meteorological forecasts, demonstrating the value of incorporating diverse data sources for urban air pollution forecasting.

Spatiotemporal modeling has emerged as a promising direction for citywide air pollution forecasting. Le et al. [2] introduced a spatiotemporal graph convolutional recurrent neural network model that effectively captures both spatial and temporal dependencies in urban air quality data, showing superior performance for city-scale predictions.

Comparative analyses of machine learning approaches have provided valuable insights into algorithm selection. Recent research [4] has conducted extensive comparative analysis and ensemble strategies for air quality forecasting, demonstrating that ensemble methods often outperform individual models through error reduction and stability improvements.

The XGBoost algorithm, introduced by Chen and Guestrin [6], has demonstrated exceptional performance in various regression tasks due to its scalable tree boosting architecture and robustness to overfitting. Its effectiveness in environmental prediction tasks has been widely recognized.

However, most existing studies focus on single-city analyses or limited pollutant types. Many approaches also suffer from inadequate feature engineering for capturing complex temporal patterns. My work distinguishes itself by simultaneously predicting six key pollutants across multiple global cities with comprehensive feature engineering, providing a more holistic approach to air quality assessment and model generalization validation.

III. MATERIALS AND METHODS

A. Dataset Description

The study utilizes a comprehensive air quality dataset containing 10,000 records from 20 major global cities, spanning the entire year of 2023. The dataset includes measurements of six primary pollutants:

- **PM2.5:** Fine particulate matter (diameter 2.5 μ m)
- **PM10:** Coarse particulate matter (diameter 10 μ m)
- **NO2:** Nitrogen dioxide
- **SO2:** Sulfur dioxide
- **CO:** Carbon monoxide
- **O3:** Ozone

Additionally, the dataset includes meteorological parameters:

- Temperature ($^{\circ}$ C)
- Relative humidity (%)
- Wind speed (m/s)

Data quality assurance involved comprehensive preprocessing steps including handling missing values, temporal alignment, and outlier detection. The AQI calculation follows the standard methodology using breakpoints for each pollutant, with the final AQI determined by the maximum score among all pollutants, capped at 500.

B. Feature Engineering

Comprehensive feature engineering was performed to capture temporal patterns and dependencies. The feature set includes:

Base Features: Original pollutant concentrations and meteorological data

Temporal Features:

- DayOfYear, Month, Year, DayOfWeek, WeekOfYear
- Quarter, IsWeekend, Season

Lag Features: AQI values from previous days (1, 2, 3, 7-day lags)

Rolling Statistics: Moving averages and standard deviations over 3 and 7-day windows

The comprehensive feature set totaling 25 features ensures that the models have sufficient information to learn the complex relationships between various factors affecting air quality.




Fig. 1. Data processing and feature engineering pipeline

C. Machine Learning Models

1) *Multi-Layer Perceptron (MLP)*: I implemented a deep neural network with three hidden layers (128-64-32 neurons) using ReLU activation and dropout regularization (0.2) to prevent overfitting. The MLP architecture is effective for learning complex non-linear relationships in the data. The model was compiled with Adam optimizer (learning rate=0.001) and trained using mean squared error loss function.

2) *Recurrent Neural Networks*: I implemented both SimpleRNN and LSTM architectures with two recurrent layers (64-32 units) and dropout regularization, specifically designed to capture temporal dependencies in time-series data. The LSTM architecture, with its gating mechanisms, is particularly suited for capturing long-term dependencies in air quality data.

3) *XGBoost*: I employed gradient boosting algorithm known for its robustness and performance on structured data. XGBoost builds an ensemble of weak prediction models in a stage-wise fashion. Separate models were trained for each pollutant, with hyperparameter optimization conducted through grid search.

D. Experimental Setup

The experimental setup was designed to ensure robust evaluation of all models. The dataset was split temporally with 80% for training and 20% for testing, preserving the temporal order of observations. Feature scaling was applied using StandardScaler to normalize all features.

For neural networks, early stopping with patience of 10 epochs was implemented to prevent overfitting. Hyperparameter tuning for XGBoost explored combinations of:

- `n_estimators`: [50, 100, 200]
- `max_depth`: [3, 6, 9]
- `learning_rate`: [0.01, 0.1, 0.2]
- `subsample`: [0.8, 0.9, 1.0]

E. Evaluation Metrics

Model performance was evaluated using multiple metrics:

- **Mean Absolute Error (MAE)**: Average magnitude of errors
- **R-squared (R^2)**: Proportion of variance explained by the model
- **Root Mean Square Error (RMSE)**: Square root of average squared errors

IV. RESULTS

A. Single-City Analysis: Mumbai

TABLE I
MODEL PERFORMANCE COMPARISON FOR POLLUTANT PREDICTION IN MUMBAI

Model	MAE	R^2	RMSE
MLP (Perceptron)	5.32	0.9501	6.82
RNN	6.99	0.9183	8.74
LSTM	4.78	0.9635	5.90
XGBoost	0.53	0.9995	0.70
XGBoost (Tuned)	0.34	0.9998	0.45

Table I presents a comprehensive comparison of model performance for pollutant prediction in Mumbai. The results demonstrate the superior performance of XGBoost models, with the tuned version achieving exceptional prediction accuracy. The significant performance gap between tree-based models and neural networks highlights the effectiveness of gradient boosting for this regression task.

The MLP model shows respectable performance with MAE of 5.32 and R^2 of 0.9501. The RNN model exhibits higher error rates, while the LSTM architecture shows improved performance with MAE of 4.78, confirming the importance of gating mechanisms for temporal data. The dramatic improvement achieved by XGBoost underscores the advantage of ensemble methods for this prediction task.




Fig. 2. Actual vs Predicted Pollutant Concentrations in Mumbai using XGBoost




Fig. 3. Actual vs Predicted AQI in Mumbai with MLP




Fig. 4. Actual vs Predicted AQI in Mumbai with RNN




Fig. 5. Actual vs Predicted AQI in Mumbai with LSTM




Fig. 6. Actual vs Predicted AQI in Mumbai with XGBoost

B. Multi-City Analysis

TABLE II
AIR QUALITY PREDICTION PERFORMANCE ACROSS CITIES USING
TUNED XGBOOST

City	Data Points	Avg AQI	MAE	R ²
Mumbai	533	206.0	0.33	0.9998
London	482	202.0	0.35	0.9998
Beijing	481	206.4	0.35	0.9998
New York	458	205.4	0.39	0.9998
Sydney	479	205.7	0.34	0.9998
Tokyo	484	207.2	0.35	0.9998

The multi-city analysis presented in Table II reveals consistent high performance across all geographical locations. The tuned XGBoost model maintains exceptional accuracy ($R^2 \geq 0.9998$) with minimal variation in MAE across cities, demonstrating remarkable generalization capability.




Fig. 7. Multi-City Air Quality Comparison and Model Performance

Figure 7 visually represents the multi-city analysis, showing the comparative performance of the tuned XGBoost model across different urban environments. The consistency in performance metrics across cities with varying average AQI values indicates that the model's predictive capability is not significantly affected by absolute pollution levels. This finding is important for practical applications, as it suggests the framework can be reliably deployed in cities with different pollution profiles.

C. Hyperparameter Tuning Impact

Hyperparameter optimization yielded significant improvements in model performance, with the tuned XGBoost model achieving a 35.8% reduction in MAE compared to the baseline XGBoost. This substantial improvement highlights the importance of parameter optimization in machine learning applications.

The optimal parameters varied across pollutants, reflecting different characteristics of each pollutant type. Learning rate emerged as the most influential hyperparameter across all pollutants, with optimal values consistently around 0.1. The subsample parameter showed optimal values between 0.8 and 0.9, indicating that some regularization through sampling is beneficial for generalization.

V. DISCUSSION AND CONCLUSION

A. Discussion

The exceptional performance of XGBoost can be attributed to several factors inherent in its design and the characteristics of air quality data. The tree-based architecture effectively learns decision boundaries for different pollution regimes, which is particularly useful given the threshold-based nature of air quality standards. The boosting mechanism sequentially corrects prediction errors, making it well-suited for complex, multi-scale patterns in environmental data.

XGBoost's ability to handle mixed data types is particularly advantageous for air quality prediction, where features include continuous measurements, categorical variables, and derived statistics. The algorithm's inherent robustness to outliers and missing values also aligns well with environmental monitoring data characteristics.

Among neural networks, LSTM demonstrated superior performance compared to MLP and SimpleRNN, confirming the

importance of capturing long-term dependencies in air quality data. However, all neural architectures showed significantly higher error rates compared to XGBoost, potentially due to limited dataset size and the effectiveness of manually engineered temporal features.

The feature importance analysis revealed that lag features consistently emerged as the most important predictors across all pollutants, highlighting strong temporal autocorrelation in air quality data. Meteorological features, especially temperature and wind speed, showed high importance for certain pollutants, reflecting the strong influence of weather conditions on pollutant dispersion.

B. Practical Applications

The achieved prediction accuracy has significant practical implications:

- **Early warning systems:** Providing advance notice of poor air quality episodes
- **Traffic management:** Informing decisions about traffic flow optimization during high-pollution periods
- **Public health advisories:** Generating accurate information for public communication
- **Urban planning:** Supporting evidence-based decisions about land use and infrastructure

The multi-city generalization capability enables rapid deployment in new urban environments with minimal retraining, making the framework valuable for cities with limited historical data.

C. Learning Outcomes

This project provided valuable learning experiences in:

- Comprehensive feature engineering for time-series data
- Comparative analysis of different machine learning architectures
- Hyperparameter optimization techniques
- Model evaluation and validation across multiple domains
- Practical implementation of environmental monitoring systems

The project demonstrated the importance of proper experimental design and rigorous evaluation in machine learning applications.

D. Conclusion

This research demonstrates the effectiveness of machine learning approaches, particularly XGBoost with comprehensive feature engineering, for accurate air quality prediction across multiple cities. The tuned XGBoost model achieved exceptional performance with MAE of 0.34 for pollutant concentration prediction and 0.85 for AQI calculation.

The multi-city analysis provides valuable insights into urban air pollution patterns and model generalization capabilities. The consistent high performance across diverse geographical locations underscores the robustness of the proposed framework. The comprehensive evaluation demonstrates that ensemble methods with proper feature engineering can achieve remarkable accuracy in environmental prediction tasks.

This work contributes to the growing body of research applying machine learning to environmental monitoring and provides a foundation for developing operational air quality forecasting systems. The proposed framework can be readily deployed for real-time air quality prediction and proactive environmental management.

REFERENCES

- [1] Z. Liu, Z. Lu, W. Zhu, J. Yuan, Z. Cao, T. Cao, S. Liu, Y. Xu, and X. Zhang, "Comparison of machine learning methods for predicting ground-level ozone pollution in Beijing," *Frontiers in Environmental Science*, vol. 13, p. 1561794, 2025.
- [2] V.-D. Le, T.-C. Bui, and S.-K. Cha, "Spatiotemporal graph convolutional recurrent neural network model for citywide air pollution forecasting," *arXiv preprint arXiv:2304.12630*, 2023.
- [3] G. Blanco, L. Barco, L. Innocenti, and C. Rossi, "Urban air pollution forecasting: A machine learning approach leveraging satellite observations and meteorological forecasts," *arXiv preprint arXiv:2405.19901*, 2024.
- [4] "Air quality forecasting using machine learning: Comparative analysis and ensemble strategies for enhanced prediction," *Water, Air, & Soil Pollution*, vol. 236, no. 1, p. 12, 2025.
- [5] G. Naresh and B. Indira, "Air pollution prediction using multivariate LSTM deep learning model," *International Journal of Intelligent Systems and Applications in Engineering*, vol. 12, no. 1, pp. 100-110, 2024.
- [6] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785-794.
- [7] World Health Organization, "Ambient air pollution: A global assessment of exposure and burden of disease," WHO Press, 2016.